

Henrique Oliveira

Santa Clara, CA | holiveira@scu.edu | +1 (847) 624-6104 | henriqueoliveira.dev | [LinkedIn](#) | [GitHub](#)

EDUCATION

Santa Clara University

Expected Graduation: December 2027

BS in Computer Science; minor in Management Information Systems

Organizations: Santa Clara Investment Fund ([Tech Research Lead](#)), Latinx Business Student Association (made [the website](#))

Coursework: Adv. Programming (C/C++), Data Structures, Embedded Systems, Operating Systems, Computer Architecture

PROFESSIONAL EXPERIENCE

PARKD | Santa Clara, CA

May 2025 – Present

Lead Software Engineer & Co-Founder

- Co-founded [Parkd](#), a parking management platform using LPR-based permit validation to deliver real-time enforcement alerts
- Organized partnerships with Santa Clara University and Bellarmine College Prep, from cold outreach to campus-wide adoption
- Built a real-time computer vision pipeline ingesting RTSP streams from 20+ cameras, running LPR inference on [edge device / GPU instance], achieving [X]% plate-read accuracy at [Y] ms p95 camera-to-alert latency
- Architected a multi-tenant backend via Node.js, Prisma, PostgreSQL on AWS with isolated data and role-based access across 6 permission tiers, validating 10k+ daily plate reads against permits and pushing violation alerts via [Redis queue / WebSockets]
- Designed React + TypeScript dashboard, informed by user interviews; featuring live violation feeds, permit search, analytics
- Built an event lakehouse on AWS S3 (Apache Iceberg) to measure LPR precision/recall used to cut false alerts by [X]%
- Orchestrated CI/CD via dev/prod branches, GitHub Actions for auto-testing, and a containerized local dev stack (Docker + Redis), enabling A/B testing of features, driving rapid feedback iteration, delivering stable deployments to 50k+ users
- Leading team of 3 student engineers; coordinating epics, weekly stand-ups, bug triage sessions, and individual async check-ins

GITHUB | San Francisco, CA

June 2026 – September 2026

Product Management Intern - AI Security Research

- Drove model evaluation strategy for LLM-based code vulnerability scanning products, benchmarking precision/recall against an 8,256-repo / 34k-CVE ground-truth corpus + existing SAST baselines to guide accurate evaluation
- Wrote 100+ KQL queries against internal data warehouse to surface backup data + build visualizations to support eval thesis
- Led 10+ AppSec customer interviews to pressure-test eval narratives, driving async customer + team alignment on epics
- Conducted competitive intel across 5 AI code-security entrants, benchmarking evals, detection approach, and bundling strategy
- Synthesized 2+ years of ambiguous internal discussions around evaluations in <2 weeks by building an AI-assisted workflow
- Authored [blog post](#) outlining evaluation methodology and LLM-based security narrative for 100k+ enterprise customers
- Prototyped RAG data layer connecting Copilot to Microsoft Fabric via REST API, enabling cross-functional teams to answer product/metrics questions in <1 min; adoption pending responsible AI checks (mainly hallucination risk)

PERMISO (acquired by [Okta](#)) | Palo Alto, CA

June 2024 – September 2024

Product Management Intern

- Collaborated with cross-functional engineering, UI/UX design, and GTM teams to align customer needs to business objectives
- Designed and validated a per-user pricing model for SaaS and cloud infrastructure integrations by running SQL queries on AWS S3, resulting in improved sales pipeline conversion, influencing \$2M in potential ARR
- Spearheaded ‘trust score’ framework tool to quantify identity security risk posture by combining identity factors with machine learning based dynamic behavioral analytics, improving how F500 customers identify high-risk identities

PROJECTS

THREADED NEXUS | MCP Server for Relationship Management

March 2026 – Present

- Developed & maintain an open-source CRM tool unifying LinkedIn, email, WhatsApp & calendar for managing connections
- Designed a fully local inference pipeline using Gemma4 13B, fine-tuned via QLoRA on my Nvidia RTX 3060 GPU with RAG based contextual memory, with remote OpenSSH access, enabling real-time contact intelligence with no 3rd party dependency
- Tech stack: Linux Server, TypeScript, Fastify, PostgreSQL, Redis, Docker, llama.cpp, QLoRA

SKILLS & INTERESTS

Human Languages: English & Portuguese (Bilingual, Native speaker), Spanish (near-Fluent, 3rd language)

Programming Languages: JavaScript/TypeScript, C/C++, Tailwind CSS, HTML5, Python, SQL/KQL/Postgres, Linux/Bash

Frontend + QA: React.js, Next.js, UI/UX Design, A/B Testing, CI/CD Orchestration, GitHub Actions, User Feedback

Backend / Infra: Node.js, AWS EC2, S3, VMware, Prisma, Docker, RestAPIs, Fastify, Database Schema Design, Data Warehousing

Interests: Playing Guitar, Music (all genres; esp. Rock, Alt, Bossa), Writing ([Substack](#)), Chess, Photography, Anything Outdoors